

STA 602 - Intro to Bayesian Statistics

Lecture 11

Li Ma

Duke University

Monte Carlo Markov Chain Diagnostics

- ▶ We have so far seen our first example of an MCMC sampler—the Gibbs sampler.
- ▶ Generally, an MCMC sampler generates a sequence of draws that are correlated

$$\boldsymbol{\theta}^{(1)}, \boldsymbol{\theta}^{(2)}, \dots$$

- ▶ We want to apply the “law of large number” on these correlated samples to get an estimate of the integral

$$E_p g = \int g(u) p(u) du$$

by

$$\frac{1}{S} \sum_{i=1}^S g(\boldsymbol{\theta}^{(i)}) \rightarrow E_p g.$$

- ▶ But when is this Monte Carlo approximation good?
- ▶ This is a harder question to answer than for standard Monte Carlo where the samples are independent. Why?

Convergence and mixing

- ▶ In order to have small error in the approximation, we need to have two things happen
 - ▶ Convergence—the chain has moved into a high probability region of the target distribution after a burn-in period, which should be discarded.
 - ▶ Good mixing—the chain should be able to move across high probability regions of the target distribution with ease, rather than getting stuck in one high probability region for a long time, and only jump to another very infrequently.

Convergence and mixing

- ▶ Convergence reduces the bias in our integral evaluation—we are in fact applying the LLN to the correct distribution p .
- ▶ Good mixing means that we give “deep” coverage on all important regions of the target distribution. Poor mixing can lead to
 - ▶ large Monte Carlo errors (assuming the chain is run long enough)
 - ▶ completely wrong estimate if the chain had barely visited some high probability regions
- ▶ How do we tell?
 - ▶ Strictly speaking, there is no rigorous way to tell with certainty if convergence and/or good mixing have been achieved. (The simple scenario like in the bivariate normal example doesn't occur in practice.)
 - ▶ We can only tell if there is clear deviation from convergence and proper mixing using heuristic strategies.
 - ▶ In particular, the chain could well be trapped around a local mode and looks as if it had converged.

Some heuristic strategies to check clear deviation

- ▶ Assessing convergence and mixing:
 - ▶ Traceplot (unbinned or binned)
 - ▶ Auto-correlation.
 - ▶ Common summary/test statistics applied to the MCMC chain.

Trace plots

- ▶ For each parameter, plot their value drawn in each iteration versus the index of iterations.
- ▶ A “good-looking” trace plot should not have
 - ▶ Drifting patterns (indicating non-convergence).
 - ▶ “S” patterns (indicating poor mixing). This can often occur as the MCMC moves are very small.
 - ▶ Infrequent jumps across regions (indicating very poor mixing).

Some examples of trace plots




Figure 1: Source: <https://jrnold.github.io>

Autocorrelation of the MCMC chain

- ▶ Autocorrelation quantifies how similar (sticky) adjacent values in the MCMC chains are.
 - ▶ Note that variables in a Markov Chain depends on the past only through the present.
 - ▶ There will still be “marginal” correlation across draws (without conditioning on the present).
 - ▶ Such auto-correlation decays with the lag.
- ▶ Lag- k autocorrelation for evaluating $E_p g = \int g(\theta)p(\theta)d\theta$:

$$\begin{aligned}\text{acf}_k(g) &= \rho_k(g) \\ &= \widehat{\text{corr}}(g^{(t)}, g^{(t+k)}) \\ &= \frac{\frac{1}{S-k} \sum_{t=1}^{S-k} (g^{(t)} - \bar{g})(g^{(t+k)} - \bar{g})}{\frac{1}{S-1} \sum_{t=1}^S (g^{(t)} - \bar{g})^2}.\end{aligned}$$

where $g^{(t)} = g(\theta^{(t)})$ and $\bar{g} = \frac{1}{S} \sum_{t=1}^S g(\theta^{(t)})$.

Effective sample size (ESS)

- ▶ The ESS is the sample size under standard Monte Carlo samples that would give the same variance as the MCMC estimate:

$$\text{Var} \bar{g} = \frac{\text{Var}_p g}{S_{eff}}.$$

- ▶ The variance of an MCMC estimate \bar{g} for $E_p g$ is given by

$$\begin{aligned}\text{Var} \bar{g} &= E [(\bar{g} - E_p g)^2] \\&= E \left[\left(\frac{1}{S} \sum_t g^{(t)} - E_p g \right)^2 \right] = \frac{1}{S^2} E \left[\sum_{s,t} (g^{(t)} - E_p g)(g^{(s)} - E_p g) \right] \\&= \frac{1}{S^2} E \left[\sum_{t=1}^S (g^{(t)} - E_p g)^2 + \sum_{t \neq s} (g^{(t)} - E_p g)(g^{(s)} - E_p g) \right] \\&= \frac{1}{S^2} \sum_{t=1}^S E [(g^{(t)} - E_p g)^2] + \frac{1}{S^2} \sum_{t \neq s} E [(g^{(t)} - E_p g)(g^{(s)} - E_p g)] \\&= \frac{1}{S} E [(g^{(t)} - E_p g)^2] + \frac{1}{S^2} \sum_{t \neq s} E [(g^{(t)} - E_p g)(g^{(s)} - E_p g)]\end{aligned}$$

Effective sample size (ESS)

- Note that in the first term

$$\frac{1}{S} \underbrace{\mathbb{E} \left[(g^{(t)} - \mathbb{E}_p g)^2 \right]}_{\text{expectation under } p \text{ when MCMC has converged}} = \frac{\text{Var}_p g}{S}.$$

because under convergence, the marginal distribution of $\theta^{(t)}$ is exactly p .

- Thus this term is exactly the variance of the sample mean (of sample size S) from a standard Monte Carlo (i.e., i.i.d. sampling).

Effective sample size (ESS)

- The second term,

$$\begin{aligned}& \frac{1}{S^2} \sum_{t \neq s} \mathbb{E} \left[(g^{(t)} - \mathbb{E}_p g)(g^{(s)} - \mathbb{E}_p g) \right] \\&= \frac{1}{S^2} \cdot 2 \sum_{k=1}^{S-1} \sum_{t=1}^{S-k} \mathbb{E} \left[(g^{(t)} - \mathbb{E}_p g)(g^{(t+k)} - \mathbb{E}_p g) \right] \\&= \frac{2}{S^2} \cdot \sum_{k=1}^{S-1} (S-k) \cdot \text{Cov}(g^{(t)}, g^{(t+k)}) \\&= \frac{2}{S^2} \cdot \sum_{k=1}^{S-1} (S-k) \cdot \text{corr}(g^{(t)}, g^{(t+k)}) \cdot \text{Var}_p g \\&\approx \frac{2}{S} \sum_{k=1}^{S-1} \text{corr}(g^{(t)}, g^{(t+k)}) \cdot \text{Var}_p g.\end{aligned}$$

- The last approximation is based on the assumption that the Lag- k autocorrelation decays sufficiently fast with k and so only for small k , the contribution is non-negligible.

Effective sample size (ESS)

- ▶ The lag- k autocorrelation $\text{corr}(g^{(t)}, g^{(t+k)})$ can be estimated by $\text{acf}_k(g)$.
- ▶ The equivalent sample size (ESS), S_{eff} , is one such that
- ▶ Thus

$$\text{Var} \bar{g} = \frac{\text{Var}_p g}{S_{eff}(g)} \quad \text{and so} \quad S_{eff}(g) = \frac{\text{Var}_p g}{\text{Var} \bar{g}}.$$

- ▶ Based on the previous calculation,

$$\begin{aligned} S_{eff}(g) &\approx \frac{\text{Var}_p g}{\frac{\text{Var}_p g}{S} + \frac{2}{S} \text{Var}_p g \cdot \sum_k \text{acf}_k(g)} \\ &\approx \frac{S}{1 + 2 \sum_k \text{acf}_k(g)}. \end{aligned}$$

- ▶ The larger the autocorrelations, the larger the variance of the MCMC estimate, and the smaller the ESS.

Back to the bivariate normal example

```
mu <- c(2,2) # mean
rho <- 0.5 # correlation - change this to different values
Sigma <- matrix(c(1,rho,rho,1),ncol=2); # covariance

S <- 1000
theta.mc <- matrix(0,nrow=S,ncol=2)
colnames(theta.mc) <- c("theta1","theta2")
theta <- c(0,0) # initial value
theta.prev <- theta

# Gibbs sampling
for (t in 1:S) {
  theta[1] <- rnorm(1,mean=mu[1]+rho*(theta[2]-mu[2]),sd=sqrt(1-rho^2))
  theta[2] <- rnorm(1,mean=mu[2]+rho*(theta[1]-mu[1]),sd=sqrt(1-rho^2))
  theta.mc[t,] <- theta
  theta.prev <- theta
}
```

For $\rho = 0.5$

```
library(coda)
theta.coda <- mcmc(theta.mc, start = 1) # no burn-in steps
options(digits=3)
summary(theta.coda)

##
## Iterations = 1:1000
## Thinning interval = 1
## Number of chains = 1
## Sample size per chain = 1000
##
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
##           Mean    SD Naive SE Time-series SE
## theta1 1.98 1.01    0.032      0.0516
## theta2 1.97 1.04    0.033      0.0437
##
## 2. Quantiles for each variable:
##
##           2.5%  25%  50%  75%  97.5%
## theta1 0.05224 1.27 2.00 2.66 3.85
## theta2 -0.00985 1.26 1.99 2.66 4.06
```

For $\rho = 0.5$

```
autocorr(theta.coda)
```

```
## , , theta1
##
##          theta1  theta2
## Lag 0    1.0000  0.5187
## Lag 1    0.2679  0.1362
## Lag 5   -0.0204 -0.0316
## Lag 10   0.0746  0.0200
## Lag 50  -0.0429 -0.0302
##
## , , theta2
##
##          theta1  theta2
## Lag 0    0.51868  1.00000
## Lag 1    0.51994  0.27462
## Lag 5   -0.00384 -0.01858
## Lag 10   0.08484  0.02438
## Lag 50  -0.01612  0.00618
```


```
effectiveSize(theta.coda)
```

```
## theta1 theta2
##    383    569
```


For $\rho = 0.5$

```
plot(theta.coda)
```


Trace of theta1




Density of theta1



Trace of theta2





Density of theta2



For $\rho = 0.5$

```
autocorr.plot(theta.coda, lag.max=100)
```



For $\rho = 0.99$

```
library(coda)
theta.coda <- mcmc(theta.mc, start = 1) # no burn-in steps
options(digits=3)
summary(theta.coda)

##
## Iterations = 1:2000
## Thinning interval = 1
## Number of chains = 1
## Sample size per chain = 2000
##
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
##           Mean      SD Naive SE Time-series SE
## theta1 1.73 0.932   0.0208         0.214
## theta2 1.73 0.930   0.0208         0.179
##
## 2. Quantiles for each variable:
##
##           2.5%  25%  50%  75%  97.5%
## theta1 -0.0518 1.12 1.75 2.37 3.36
## theta2 -0.1001 1.11 1.74 2.37 3.34
```

For $\rho = 0.99$

```
autocorr(theta.coda)
```

```
## , , theta1
##
##          theta1 theta2
## Lag 0    1.000  0.989
## Lag 1    0.977  0.965
## Lag 5    0.885  0.874
## Lag 10   0.770  0.762
## Lag 50   0.313  0.309
##
## , , theta2
##
##          theta1 theta2
## Lag 0    0.989  1.000
## Lag 1    0.988  0.977
## Lag 5    0.896  0.884
## Lag 10   0.781  0.772
## Lag 50   0.318  0.315
```


```
effectiveSize(theta.coda)
```

```
## theta1 theta2
##    19.0   27.1
```


For $\rho = 0.99$

```
plot(theta.coda)
```


Trace of theta1




Density of theta1



Trace of theta2





Density of theta2



For $\rho = 0.99$

```
autocorr.plot(theta.coda, lag.max=100)
```



Thinning the chain

- ▶ Sometimes when we run the MCMC very long, there might be too many iterations that need to be saved.
- ▶ Can we save just a subset of the chain with the least impact on the resulting Monte Carlo error?
- ▶ Save every few iterations rather than every iteration.
- ▶ The stronger the auto-correlation of the chain, the more steps we can skip to maintain an efficient sample size.

Example: Air pollutant measurements

```
x <- c(104,105,103,102,105,107,106,104,103,106) # the data
n <- length(x) # sample size

S <- 10000
xbar <- mean(x)
s2 <- var(x)
n <- length(x)

THETA <- matrix(NA,nrow=S,ncol=2,dimnames=list(1:S,c("theta","sigma2")))
THETA.init <- c(xbar,s2) # Initial values set to the MLE
THETA.curr <- THETA.init # the parameter values at the current iteration
mu.0 <- 100; nu.0 <- 1; tau2.0 <- 25; sigma2.0 <- 4


### Start Gibbs sampling
for (t in 1:S) {
  tau2.n <- 1/(1/tau2.0 + n/THETA.curr[2])
  mu.n <- (mu.0/tau2.0 + xbar*n/THETA.curr[2])/(1/tau2.0 + n/THETA.curr[2])

  ## Update theta
  THETA.curr[1] <- rnorm(1,mean=mu.n,sd=sqrt(tau2.n))


  ## Update sigma2
  THETA.curr[2] <- 1/rgamma(1,shape=(nu.0+n)/2,
                           rate=1/2*(nu.0*sigma2.0+sum((x-THETA.curr[1])^2)))

  ## Save the current iteration
  THETA[t,] <- THETA.curr
}
```


Trace plot for θ



Histogram for θ




Trace plot for σ^2



Histogram for σ^2

Histogram of sampled σ^2



MCMC diagnostics

```
library(coda)
THETA.coda <- mcmc(THETA, start = 1) # no burn-in steps
options(digits=3)
summary(THETA.coda)
```

```
##
## Iterations = 1:10000
## Thinning interval = 1
## Number of chains = 1
## Sample size per chain = 10000
##
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
##           Mean      SD Naive SE Time-series SE
## theta  104.4 0.575  0.00575      0.00575
## sigma2   3.3 1.848  0.01848      0.02020
##
## 2. Quantiles for each variable:
##
##           2.5%    25%    50%    75%   97.5%
## theta  103.28 104.08 104.44 104.80 105.57
## sigma2   1.31   2.11   2.85   3.94   8.02
```

Autocorrelation and ESS

```
autocorr(THETA.coda)
```

```
## , , theta
##
##          theta    sigma2
## Lag 0    1.00000 -0.05283
## Lag 1   -0.00154  0.01328
## Lag 5    0.00377  0.00721
## Lag 10   0.00825 -0.00519
## Lag 50  -0.01892 -0.01137
##
## , , sigma2
##
##          theta    sigma2
## Lag 0   -0.05283  1.00e+00
## Lag 1   -0.05606  1.11e-01
## Lag 5    0.01465  3.28e-05
## Lag 10  -0.02335 -8.76e-03
## Lag 50  -0.00646 -1.14e-02
```


```
effectiveSize(THETA.coda)
```

```
## theta sigma2
## 10000  8372
```


Trace plots

```
plot (THETA.coda)
```


Trace of theta




Density of theta



Trace of sigma2





Density of sigma2



Autocorrelation plots

```
autocorr.plot(THETA.coda, lag.max=100)
```



Using multiple chains for MCMC diagnostics

- ▶ Using multiple chains to examine convergence and mixing is a common technique.
- ▶ If an MCMC chain has been run long enough after iteration S (i.e., converged and mixed well), then if we run M separate chains with *overdispersed* initial values, then after iteration S , they should all
 - ▶ have converged (i.e., moved to a high-probability region);
 - ▶ be able to cover all high-probability regions sufficiently deep.
- ▶ These imply that the distributions of draws from the M chains should all be similar.
- ▶ So we can construct various statistics to examine their differences.

Gelman-Rubin statistic

- ▶ Start off from running M separate chains with overdispersed starting points (with respect to the target p) for $2S$ iterations.
- ▶ Discard the first S draws in each as burn-in.
- ▶ Suppose the remaining S iterations are

$$\begin{array}{c} g_1^{(1)}, g_1^{(2)}, \dots, g_1^{(S)} \\ g_2^{(1)}, g_2^{(2)}, \dots, g_2^{(S)} \\ \vdots \\ g_M^{(1)}, g_M^{(2)}, \dots, g_M^{(S)} \end{array}$$

- ▶ Let's use these $S \cdot M$ draws to estimate $\text{Var}_p g$.

Main idea

- ▶ If convergence is achieved and mixing is fine, then an unbiased estimate for $\text{Var} \bar{g}$ should be similar whether or not we use draws from within each chain, and those from across chains.
- ▶ Otherwise, we will overestimate $\text{Var}_p g$ if we use draws across chains and underestimate if we only use draws from within each chain.
- ▶ Strategy:
 - ▶ Compute the ratio \hat{R} of these two estimate (and take a square root).
 - ▶ Examine how close \hat{R} is to 1.

The Gelman-Rubin statistics

- Specifically, the Gelman-Rubin statistic is

$$\hat{R} = \sqrt{\frac{\frac{S-1}{S}W + \frac{1}{S}B}{W}}$$

where

$$B = \frac{S}{M-1} \sum_{m=1}^M (\bar{g}_m - \bar{g}_{\cdot})^2 \quad \text{and} \quad W = \frac{1}{M} \sum_{m=1}^M s_m(g)^2$$

where $\bar{g}_m = \frac{1}{S} \sum_t g_m^{(t)}$, $\bar{g}_{\cdot} = \frac{1}{M} \sum_{m=1}^M \bar{g}_m$, and

$$s_m(g)^2 = \frac{1}{S-1} \sum_{t=1}^S (g_m^{(t)} - \bar{g}_m)^2.$$

- As $S \rightarrow \infty$, $\hat{R} \rightarrow 1$.
- Gelman (2004) recommends running the chains until $\hat{R} < 1.1$, though this is just a rule-of-thumb based on experience. There is not much theoretical support for that.

Using different portions of a single chain

- ▶ Another strategy for examining convergence and mixing is by comparing disjoint portions of a single chain.
 - ▶ Geweke (1992) proposes to compare the MCMC estimates based on two disjoint portions of a single chain, and quantify the difference using a z-score.
 - ▶ Raftery and Lewis diagnostic focuses on variability in the quantiles.
- ▶ Strategies using a single chain is less robust as it can never tell if a chain is simply stuck in a high-probability region!
- ▶ Using multiple chains based on overdispersed starting points has a chance of detecting that.
- ▶ Better yet: A variant of Gelman-Rubin combines the two strategies, by using M separate chains of length S and divide each into d bins, resulting in a total of dM chains of length S/d (often $d = 2$).

Bivariate example

```
M <- 10; S <- 1000;
mu=c(2,2)
rho <- 0.99
theta.mc <- matrix(NA,nrow=S,ncol=2)
theta.mcmc.multiple <- mcmc.list()
theta.init.multiple <- mvtnorm::rmvnorm(M,mean=mu, sigma=10*matrix(c(1,0,0,1),ncol=2))
for (m in 1:M) {
  # Begin Gibbs sampling
  theta <- theta.init.multiple[m,] # initial value
  theta.prev <- theta

  for (t in 1:S) {
    theta[1] <- rnorm(1,mean=mu[1]+rho*(theta[2]-mu[2]),sd=sqrt(1-rho^2))
    theta[2] <- rnorm(1,mean=mu[2]+rho*(theta[1]-mu[1]),sd=sqrt(1-rho^2))
    theta.mc[t,] <- theta
    theta.prev <- theta
  }
  theta.coda = mcmc(theta.mc,start=1) # no burnin steps
  theta.mcmc.multiple[[m]] = theta.coda
}
gelman.diag(theta.mcmc.multiple)
```

```
## Potential scale reduction factors:
##
##      Point est. Upper C.I.
## [1,]      1.07      1.14
## [2,]      1.07      1.14
##
## Multivariate psrf
##
## 1.07
```

Bivariate example

```
gelman.plot(theta.mcmc.multiple)
```

